

Project 1 - Real Estate

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Project Task: Week 1

Data Import and Preparation:

1. Import data.
2. Figure out the primary key and look for the requirement of indexing.
3. Gauge the fill rate of the variables and devise plans for missing value treatment. Please explain explicitly the reason for the treatment chosen for each variable.

Exploratory Data Analysis (EDA):

4. Perform debt analysis. You may take the following steps:

1. Import Data

```
In [2]: df_train = pd.read_csv('p1train.csv')
```

```
In [3]: df_test = pd.read_csv('p1test.csv')
```

```
In [4]: df_train.columns
```

```
Out[4]: Index(['UID', 'BLOCKID', 'SUMLEVEL', 'COUNTYID', 'STATEID', 'state',  
              'state_ab', 'city', 'place', 'type', 'primary', 'zip_code', 'area_code',  
              'lat', 'lng', 'ALand', 'AWater', 'pop', 'male_pop', 'female_pop',  
              'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight',  
              'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25',  
              'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50',  
              'universe_samples', 'used_samples', 'hi_mean', 'hi_median', 'hi_stdev',  
              'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median',  
              'family_stdev', 'family_sample_weight', 'family_samples',  
              'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',  
              'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean',  
              'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',  
              'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',  
              'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',  
              'hs_degree_male', 'hs_degree_female', 'male_age_mean',  
              'male_age_median', 'male_age_stdev', 'male_age_sample_weight',  
              'male_age_samples', 'female_age_mean', 'female_age_median',  
              'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',  
              'pct_own', 'married', 'married_snp', 'separated', 'divorced'],  
             dtype='object')
```

```
In [5]: df_test.columns
```

```
Out[5]: Index(['UID', 'BLOCKID', 'SUMLEVEL', 'COUNTYID', 'STATEID', 'state',  
              'state_ab', 'city', 'place', 'type', 'primary', 'zip_code', 'area_code',  
              'lat', 'lng', 'ALand', 'AWater', 'pop', 'male_pop', 'female_pop',  
              'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight',  
              'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25',  
              'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50',  
              'universe_samples', 'used_samples', 'hi_mean', 'hi_median', 'hi_stdev',  
              'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median',  
              'family_stdev', 'family_sample_weight', 'family_samples',  
              'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',  
              'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean',  
              'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',  
              'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',  
              'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',  
              'hs_degree_male', 'hs_degree_female', 'male_age_mean',  
              'male_age_median', 'male_age_stdev', 'male_age_sample_weight',  
              'male_age_samples', 'female_age_mean', 'female_age_median',  
              'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',  
              'pct_own', 'married', 'married_snp', 'separated', 'divorced'],  
              dtype='object')
```

```
In [6]: len(df_train)
```

```
Out[6]: 27321
```

```
In [7]: len(df_test)
```

```
Out[7]: 11709
```

In [8]: df_train.head()

Out[8]:

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	...	female_age_mean	female_age_median
0	267822	NaN	140	53	36	New York	NY	Hamilton	Hamilton	City	...	44.48629	45.33333
1	246444	NaN	140	141	18	Indiana	IN	South Bend	Roseland	City	...	36.48391	37.58333
2	245683	NaN	140	63	18	Indiana	IN	Danville	Danville	City	...	42.15810	42.83333
3	279653	NaN	140	127	72	Puerto Rico	PR	San Juan	Guaynabo	Urban	...	47.77526	50.58333
4	247218	NaN	140	161	20	Kansas	KS	Manhattan	Manhattan City	City	...	24.17693	21.58333

5 rows × 80 columns



In [9]: df_test.head()

Out[9]:

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	...	female_age_mean	female_age_r
0	255504	NaN	140	163	26	Michigan	MI	Detroit	Dearborn Heights City	CDP	...	34.78682	33
1	252676	NaN	140	1	23	Maine	ME	Auburn	Auburn City	City	...	44.23451	46
2	276314	NaN	140	15	42	Pennsylvania	PA	Pine City	Millerton	Borough	...	41.62426	44
3	248614	NaN	140	231	21	Kentucky	KY	Monticello	Monticello City	City	...	44.81200	48
4	286865	NaN	140	355	48	Texas	TX	Corpus Christi	Edroy	Town	...	40.66618	42

5 rows × 80 columns




```
In [10]: df_train.describe()
```

```
Out[10]:
```

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	zip_code	area_code	lat	lng	ALand
count	27321.000000	0.0	27321.0	27321.000000	27321.000000	27321.000000	27321.000000	27321.000000	27321.000000	2.732100e+04
mean	257331.996303	NaN	140.0	85.646426	28.271806	50081.999524	596.507668	37.508813	-91.288394	1.295106e+08
std	21343.859725	NaN	0.0	98.333097	16.392846	29558.115660	232.497482	5.588268	16.343816	1.275531e+09
min	220342.000000	NaN	140.0	1.000000	1.000000	602.000000	201.000000	17.929085	-165.453872	4.113400e+04
25%	238816.000000	NaN	140.0	29.000000	13.000000	26554.000000	405.000000	33.899064	-97.816067	1.799408e+06
50%	257220.000000	NaN	140.0	63.000000	28.000000	47715.000000	614.000000	38.755183	-86.554374	4.866940e+06
75%	275818.000000	NaN	140.0	109.000000	42.000000	77093.000000	801.000000	41.380606	-79.782503	3.359820e+07
max	294334.000000	NaN	140.0	840.000000	72.000000	99925.000000	989.000000	67.074017	-65.379332	1.039510e+11

8 rows × 74 columns



```
In [11]: df_train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27321 entries, 0 to 27320
Data columns (total 80 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   UID                 27321 non-null  int64
 1   BLOCKID             0 non-null      float64
 2   SUMLEVEL            27321 non-null  int64
 3   COUNTYID            27321 non-null  int64
 4   STATEID             27321 non-null  int64
 5   state               27321 non-null  object
 6   state_ab            27321 non-null  object
 7   city                27321 non-null  object
 8   place               27321 non-null  object
 9   type                27321 non-null  object
10  primary              27321 non-null  object
11  zip_code             27321 non-null  int64
12  area_code            27321 non-null  int64
13  lat                  27321 non-null  float64
14  lng                  27321 non-null  float64
15  ALand                27321 non-null  float64
16  AWater               27321 non-null  int64
17  pop                  27321 non-null  int64
18  male_pop             27321 non-null  int64
19  female_pop           27321 non-null  int64
20  rent_mean            27007 non-null  float64
21  rent_median          27007 non-null  float64
22  rent_stdev           27007 non-null  float64
23  rent_sample_weight   27007 non-null  float64
24  rent_samples         27007 non-null  float64
25  rent_gt_10           27007 non-null  float64
26  rent_gt_15           27007 non-null  float64
27  rent_gt_20           27007 non-null  float64
28  rent_gt_25           27007 non-null  float64
29  rent_gt_30           27007 non-null  float64
30  rent_gt_35           27007 non-null  float64
31  rent_gt_40           27007 non-null  float64
32  rent_gt_50           27007 non-null  float64
33  universe_samples     27321 non-null  int64
34  used_samples         27321 non-null  int64
35  hi_mean              27053 non-null  float64

```

36	hi_median	27053	non-null	float64
37	hi_stdev	27053	non-null	float64
38	hi_sample_weight	27053	non-null	float64
39	hi_samples	27053	non-null	float64
40	family_mean	27023	non-null	float64
41	family_median	27023	non-null	float64
42	family_stdev	27023	non-null	float64
43	family_sample_weight	27023	non-null	float64
44	family_samples	27023	non-null	float64
45	hc_mortgage_mean	26748	non-null	float64
46	hc_mortgage_median	26748	non-null	float64
47	hc_mortgage_stdev	26748	non-null	float64
48	hc_mortgage_sample_weight	26748	non-null	float64
49	hc_mortgage_samples	26748	non-null	float64
50	hc_mean	26721	non-null	float64
51	hc_median	26721	non-null	float64
52	hc_stdev	26721	non-null	float64
53	hc_samples	26721	non-null	float64
54	hc_sample_weight	26721	non-null	float64
55	home_equity_second_mortgage	26864	non-null	float64
56	second_mortgage	26864	non-null	float64
57	home_equity	26864	non-null	float64
58	debt	26864	non-null	float64
59	second_mortgage_cdf	26864	non-null	float64
60	home_equity_cdf	26864	non-null	float64
61	debt_cdf	26864	non-null	float64
62	hs_degree	27131	non-null	float64
63	hs_degree_male	27121	non-null	float64
64	hs_degree_female	27098	non-null	float64
65	male_age_mean	27132	non-null	float64
66	male_age_median	27132	non-null	float64
67	male_age_stdev	27132	non-null	float64
68	male_age_sample_weight	27132	non-null	float64
69	male_age_samples	27132	non-null	float64
70	female_age_mean	27115	non-null	float64
71	female_age_median	27115	non-null	float64
72	female_age_stdev	27115	non-null	float64
73	female_age_sample_weight	27115	non-null	float64
74	female_age_samples	27115	non-null	float64
75	pct_own	27053	non-null	float64
76	married	27130	non-null	float64
77	married_snp	27130	non-null	float64


```
78  separated      27130 non-null  float64
79  divorced      27130 non-null  float64
dtypes: float64(62), int64(12), object(6)
memory usage: 16.7+ MB
```

In [12]: `df_test.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 11709 entries, 0 to 11708
```

```
Data columns (total 80 columns):
```

#	Column	Non-Null Count	Dtype
0	UID	11709 non-null	int64
1	BLOCKID	0 non-null	float64
2	SUMLEVEL	11709 non-null	int64
3	COUNTYID	11709 non-null	int64
4	STATEID	11709 non-null	int64
5	state	11709 non-null	object
6	state_ab	11709 non-null	object
7	city	11709 non-null	object
8	place	11709 non-null	object
9	type	11709 non-null	object
10	primary	11709 non-null	object
11	zip_code	11709 non-null	int64
12	area_code	11709 non-null	int64
13	lat	11709 non-null	float64
14	lng	11709 non-null	float64
15	ALand	11709 non-null	int64
16	AWater	11709 non-null	int64
17	pop	11709 non-null	int64
18	male_pop	11709 non-null	int64
19	female_pop	11709 non-null	int64
20	rent_mean	11561 non-null	float64
21	rent_median	11561 non-null	float64
22	rent_stdev	11561 non-null	float64
23	rent_sample_weight	11561 non-null	float64
24	rent_samples	11561 non-null	float64
25	rent_gt_10	11560 non-null	float64
26	rent_gt_15	11560 non-null	float64
27	rent_gt_20	11560 non-null	float64
28	rent_gt_25	11560 non-null	float64
29	rent_gt_30	11560 non-null	float64
30	rent_gt_35	11560 non-null	float64
31	rent_gt_40	11560 non-null	float64
32	rent_gt_50	11560 non-null	float64
33	universe_samples	11709 non-null	int64
34	used_samples	11709 non-null	int64
35	hi_mean	11587 non-null	float64

36	hi_median	11587	non-null	float64
37	hi_stdev	11587	non-null	float64
38	hi_sample_weight	11587	non-null	float64
39	hi_samples	11587	non-null	float64
40	family_mean	11573	non-null	float64
41	family_median	11573	non-null	float64
42	family_stdev	11573	non-null	float64
43	family_sample_weight	11573	non-null	float64
44	family_samples	11573	non-null	float64
45	hc_mortgage_mean	11441	non-null	float64
46	hc_mortgage_median	11441	non-null	float64
47	hc_mortgage_stdev	11441	non-null	float64
48	hc_mortgage_sample_weight	11441	non-null	float64
49	hc_mortgage_samples	11441	non-null	float64
50	hc_mean	11419	non-null	float64
51	hc_median	11419	non-null	float64
52	hc_stdev	11419	non-null	float64
53	hc_samples	11419	non-null	float64
54	hc_sample_weight	11419	non-null	float64
55	home_equity_second_mortgage	11489	non-null	float64
56	second_mortgage	11489	non-null	float64
57	home_equity	11489	non-null	float64
58	debt	11489	non-null	float64
59	second_mortgage_cdf	11489	non-null	float64
60	home_equity_cdf	11489	non-null	float64
61	debt_cdf	11489	non-null	float64
62	hs_degree	11624	non-null	float64
63	hs_degree_male	11620	non-null	float64
64	hs_degree_female	11604	non-null	float64
65	male_age_mean	11625	non-null	float64
66	male_age_median	11625	non-null	float64
67	male_age_stdev	11625	non-null	float64
68	male_age_sample_weight	11625	non-null	float64
69	male_age_samples	11625	non-null	float64
70	female_age_mean	11613	non-null	float64
71	female_age_median	11613	non-null	float64
72	female_age_stdev	11613	non-null	float64
73	female_age_sample_weight	11613	non-null	float64
74	female_age_samples	11613	non-null	float64
75	pct_own	11587	non-null	float64
76	married	11625	non-null	float64
77	married_snp	11625	non-null	float64

```

78  separated      11625 non-null  float64
79  divorced      11625 non-null  float64
dtypes: float64(61), int64(13), object(6)
memory usage: 7.1+ MB

```

2. Figure out the primary key and look for the requirement of indexing

```

In [13]: df_train.set_index(keys=['UID'],inplace=True)
df_test.set_index(keys=['UID'],inplace=True)

```

```


In [14]: df_train.head()

```

Out[14]:

	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	primary	...	female_age_mean	female_age_n
UID													
267822	NaN	140	53	36	New York	NY	Hamilton	Hamilton	City	tract	...	44.48629	45
246444	NaN	140	141	18	Indiana	IN	South Bend	Roseland	City	tract	...	36.48391	37
245683	NaN	140	63	18	Indiana	IN	Danville	Danville	City	tract	...	42.15810	42
279653	NaN	140	127	72	Puerto Rico	PR	San Juan	Guaynabo	Urban	tract	...	47.77526	50
247218	NaN	140	161	20	Kansas	KS	Manhattan	Manhattan City	City	tract	...	24.17693	21

5 rows × 79 columns



```
In [15]: df_test.head()
```

```
Out[15]:
```

	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	primary	...	female_age_mean	female
UID													
255504	NaN	140	163	26	Michigan	MI	Detroit	Dearborn Heights City	CDP	tract	...	34.78682	
252676	NaN	140	1	23	Maine	ME	Auburn	Auburn City	City	tract	...	44.23451	
276314	NaN	140	15	42	Pennsylvania	PA	Pine City	Millerton	Borough	tract	...	41.62426	
248614	NaN	140	231	21	Kentucky	KY	Monticello	Monticello City	City	tract	...	44.81200	
286865	NaN	140	355	48	Texas	TX	Corpus Christi	Edroy	Town	tract	...	40.66618	

5 rows × 79 columns



3. Gauge the fill rate of the variables and devise plans for missing value treatment. Please explain explicitly the reason for the treatment chosen for each variable.

```
In [16]: missing_list_train = df_train.isnull().sum() * 100/len(df_train)
```

```
In [17]: missing_values_df_train = pd.DataFrame(missing_list_train,columns=['Percentage of Missing Values'])
```

```
In [18]: missing_values_df_train.sort_values(by='Percentage of Missing Values',inplace=True,ascending=False)
```

```
In [19]: missing_values_df_train[missing_values_df_train['Percentage of Missing Values'] > 0][:10]
```

Out[19]:

Percentage of Missing Values	
BLOCKID	100.000000
hc_samples	2.196113
hc_mean	2.196113
hc_median	2.196113
hc_stdev	2.196113
hc_sample_weight	2.196113
hc_mortgage_mean	2.097288
hc_mortgage_stdev	2.097288
hc_mortgage_sample_weight	2.097288
hc_mortgage_samples	2.097288

```
In [20]: missing_list_test = df_test.isnull().sum() *100/len(df_train)
```

```
In [21]: missing_values_df_test = pd.DataFrame(missing_list_test,columns=['Percentage of Missing Values'])
```

```
In [22]: missing_values_df_test.sort_values(by='Percentage of Missing Values',inplace=True,ascending=False)
```

```
In [23]: missing_values_df_test[missing_values_df_test['Percentage of Missing Values']>0][:10]
```

Out[23]:

Percentage of Missing Values	
BLOCKID	42.857143
hc_samples	1.061455
hc_mean	1.061455
hc_median	1.061455
hc_stdev	1.061455
hc_sample_weight	1.061455
hc_mortgage_mean	0.980930
hc_mortgage_stdev	0.980930
hc_mortgage_sample_weight	0.980930
hc_mortgage_samples	0.980930

```
In [24]: df_train.drop(columns=['BLOCKID', 'SUMLEVEL'], inplace=True)
# BLOCKID can be dropped, since it has 100% missing Values
# SUMLEVEL can be dropped, since it does not have any predictive power and no variance
```

```
In [25]: df_test.drop(columns=['BLOCKID', 'SUMLEVEL'], inplace=True)
# BLOCKID can be dropped, since it has 43% missing Values
# SUMLEVEL can be dropped, since it does not have any predictive power and no variance
```


In [26]: *# Imputing the missing values in other columns with mean*

```
missing_train_cols = []
for col in df_train.columns:
    if df_train[col].isna().sum() != 0:
        missing_train_cols.append(col)
print(missing_train_cols)
```

```
['rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_g
t_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50', 'hi_mean', 'hi_median', 'hi_stdev', 'hi
_sample_weight', 'hi_samples', 'family_mean', 'family_median', 'family_stdev', 'family_sample_weight', 'family_sampl
es', 'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_sample
s', 'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight', 'home_equity_second_mortgage', 'second_mor
tgage', 'home_equity', 'debt', 'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree', 'hs_degree_male',
'hs_degree_female', 'male_age_mean', 'male_age_median', 'male_age_stdev', 'male_age_sample_weight', 'male_age_sample
s', 'female_age_mean', 'female_age_median', 'female_age_stdev', 'female_age_sample_weight', 'female_age_samples', 'p
ct_own', 'married', 'married_snp', 'separated', 'divorced']
```

In [27]:

```
missing_test_cols = []
for col in df_test.columns:
    if df_test[col].isna().sum() != 0:
        missing_test_cols.append(col)
print(missing_test_cols)
```

```
['rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_g
t_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50', 'hi_mean', 'hi_median', 'hi_stdev', 'hi
_sample_weight', 'hi_samples', 'family_mean', 'family_median', 'family_stdev', 'family_sample_weight', 'family_sampl
es', 'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_sample
s', 'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight', 'home_equity_second_mortgage', 'second_mor
tgage', 'home_equity', 'debt', 'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree', 'hs_degree_male',
'hs_degree_female', 'male_age_mean', 'male_age_median', 'male_age_stdev', 'male_age_sample_weight', 'male_age_sample
s', 'female_age_mean', 'female_age_median', 'female_age_stdev', 'female_age_sample_weight', 'female_age_samples', 'p
ct_own', 'married', 'married_snp', 'separated', 'divorced']
```

In [28]:

```
for col in df_train.columns:
    if col in (missing_train_cols):
        df_train[col].replace(np.nan, df_train[col].mean(), inplace=True)
```

```
In [29]: for col in df_test.columns:
         if col in (missing_test_cols):
             df_test[col].replace(np.nan, df_test[col].mean(),inplace=True)
```

```
In [30]: df_train.isnull().sum().sum()
```

```
Out[30]: 0
```

```
In [31]: df_test.isnull().sum().sum()
```

```
Out[31]: 0
```

Exploratory Data Analysis (EDA):

4.Perform debt analysis. You may take the following steps:

- a) Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map. You may keep the upper limit for the percent of households with a second mortgage to 50 percent
 - b) Use the following bad debt equation: $\text{Bad Debt} = P (\text{Second Mortgage} \times \text{Home Equity Loan})$
 $\text{Bad Debt} = \text{second_mortgage} + \text{home_equity} - \text{home_equity_second_mortgage}$
 - c) Create pie charts to show overall debt and bad debt
 - d) Create Box and whisker plot and analyze the distribution for 2nd mortgage, home equity, good debt, and bad debt for different cities
 - e) Create a collated income distribution chart for family income, house hold income, and remaining income
-
- a) Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map. You may keep the upper limit for the percent of households with a second mortgage to 50 percent

```
In [32]: from pandasql import sqldf
```

```
In [33]: q1 = "select place,pct_own,second_mortgage,lat,lng from df_train where pct_own > 0.10 and second_mortgage < 0.5 order
```

```
In [34]: pysqldf = lambda q: sqldf(q, globals())
```

```
In [35]: df_train_location_mort_pct = pysqldf(q1)
```

```
In [36]: df_train_location_mort_pct.head()
```

Out[36]:

	place	pct_own	second_mortgage	lat	lng
0	Worcester City	0.20247	0.43363	42.254262	-71.800347
1	Harbor Hills	0.15618	0.31818	40.751809	-73.853582
2	Glen Burnie	0.22380	0.30212	39.127273	-76.635265
3	Egypt Lake-leto	0.11618	0.28972	28.029063	-82.495395
4	Lincolnwood	0.14228	0.28899	41.967289	-87.652434


```
In [37]: import plotly.express as px
```

```
In [38]: import plotly.graph_objects as go
```

```
In [39]: fig = go.Figure(data=go.Scattergeo(
    lat = df_train_location_mort_pct['lat'],
    lon = df_train_location_mort_pct['lng'],
    ))
```

```
In [40]: fig.update_layout(  
    geo=dict(  
        scope = 'north america',  
        showland = True,  
        landcolor = "rgb(212, 212, 212)",  
        subunitcolor = "rgb(255, 255, 255)",  
        countrycolor = "rgb(255, 255, 255)",  
        showlakes = True,  
        lakecolor = "rgb(255, 255, 255)",  
        showsubunits = True,  
        showcountries = True,  
        resolution = 50,  
        projection = dict(  
            type = 'conic conformal',  
            rotation_lon = -100  
        ),  
        lonaxis = dict(  
            showgrid = True,  
            gridwidth = 0.5,  
            range = [-140.0, -55.0],  
            dtick = 5  
        ),  
        lataxis = dict(  
            showgrid = True,  
            gridwidth = 0.5,  
            range = [20.0, 60.0],  
            dtick = 5  
        )  
    ),  
    title = 'Top 2500 locations with second mortgage is the highest and percent ownership is above 10 percent')  
fig.show()
```

Top 2500 locations with second mortgage is the highest and percent ownership is above 10 percent



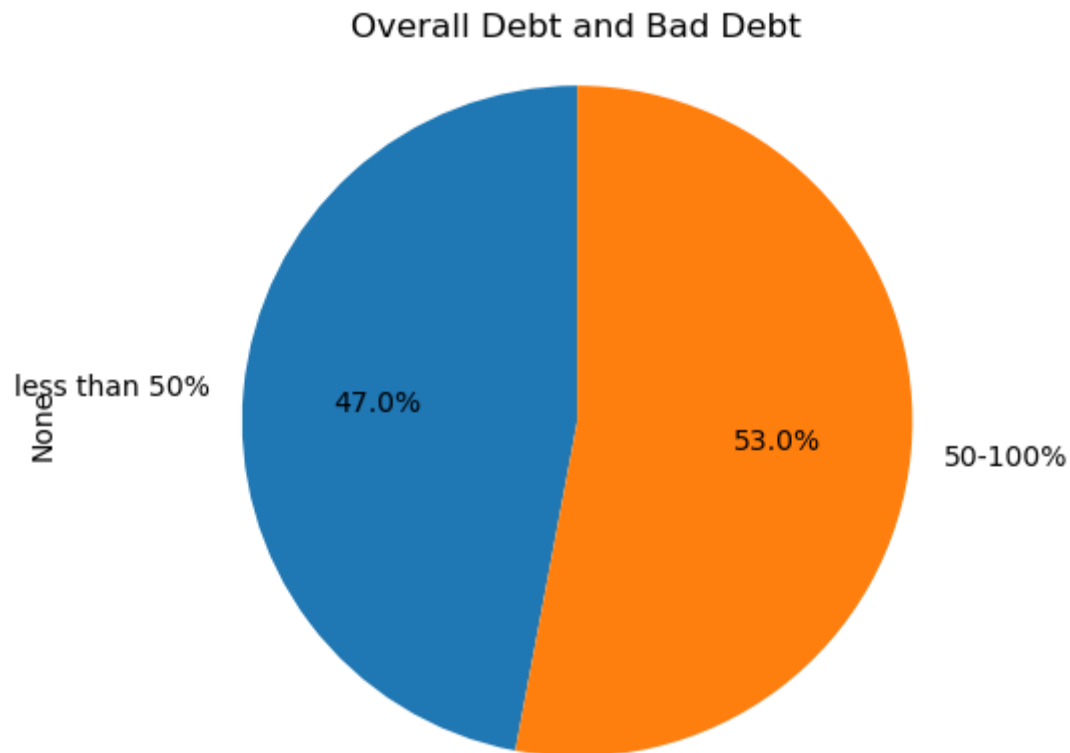
b) Use the following bad debt equation: $\text{Bad Debt} = P(\text{Second Mortgage} \cap \text{Home Equity Loan})$ $\text{Bad Debt} = \text{second_mortgage} + \text{home_equity} - \text{home_equity_second_mortgage}$

```
In [41]: df_train['bad_debt'] = df_train['second_mortgage'] + df_train['home_equity'] - df_train['home_equity_second_mortgage']
```

c) Create pie charts to show overall debt and bad debt

```
In [42]: df_train['bins'] = pd.cut(df_train['bad_debt'],bins=[0,0.10,1],labels=['less than 50%','50-100%'])
```

```
In [43]: df_train.groupby(['bins']).size().plot(kind='pie',subplots=True,startangle=90,autopct='%1.1f%%')  
plt.axis('equal')  
plt.title('Overall Debt and Bad Debt')  
plt.show()
```



d) Create Box and whisker plot and analyze the distribution for 2nd mortgage, home equity, good debt, and bad debt for different cities

```
In [44]: cols = []  
df_train.columns
```

```
Out[44]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',  
              'primary', 'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater',  
              'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',  
              'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',  
              'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',  
              'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',  
              'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',  
              'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',  
              'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',  
              'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',  
              'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',  
              'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',  
              'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',  
              'hs_degree_male', 'hs_degree_female', 'male_age_mean',  
              'male_age_median', 'male_age_stdev', 'male_age_sample_weight',  
              'male_age_samples', 'female_age_mean', 'female_age_median',  
              'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',  
              'pct_own', 'married', 'married_snp', 'separated', 'divorced',  
              'bad_debt', 'bins'],  
             dtype='object')
```

```
In [45]: cols = ['second mortgage', 'home_equity', 'debt', 'bad_dept']  
df_box_hamilton = df_train.loc[df_train.city == 'Hamilton']  
df_box_manhattan = df_train.loc[df_train.city == 'Manhattan']  
df_box_city = pd.concat([df_box_hamilton, df_box_manhattan])
```

```
In [46]: df_box_city.head()
```

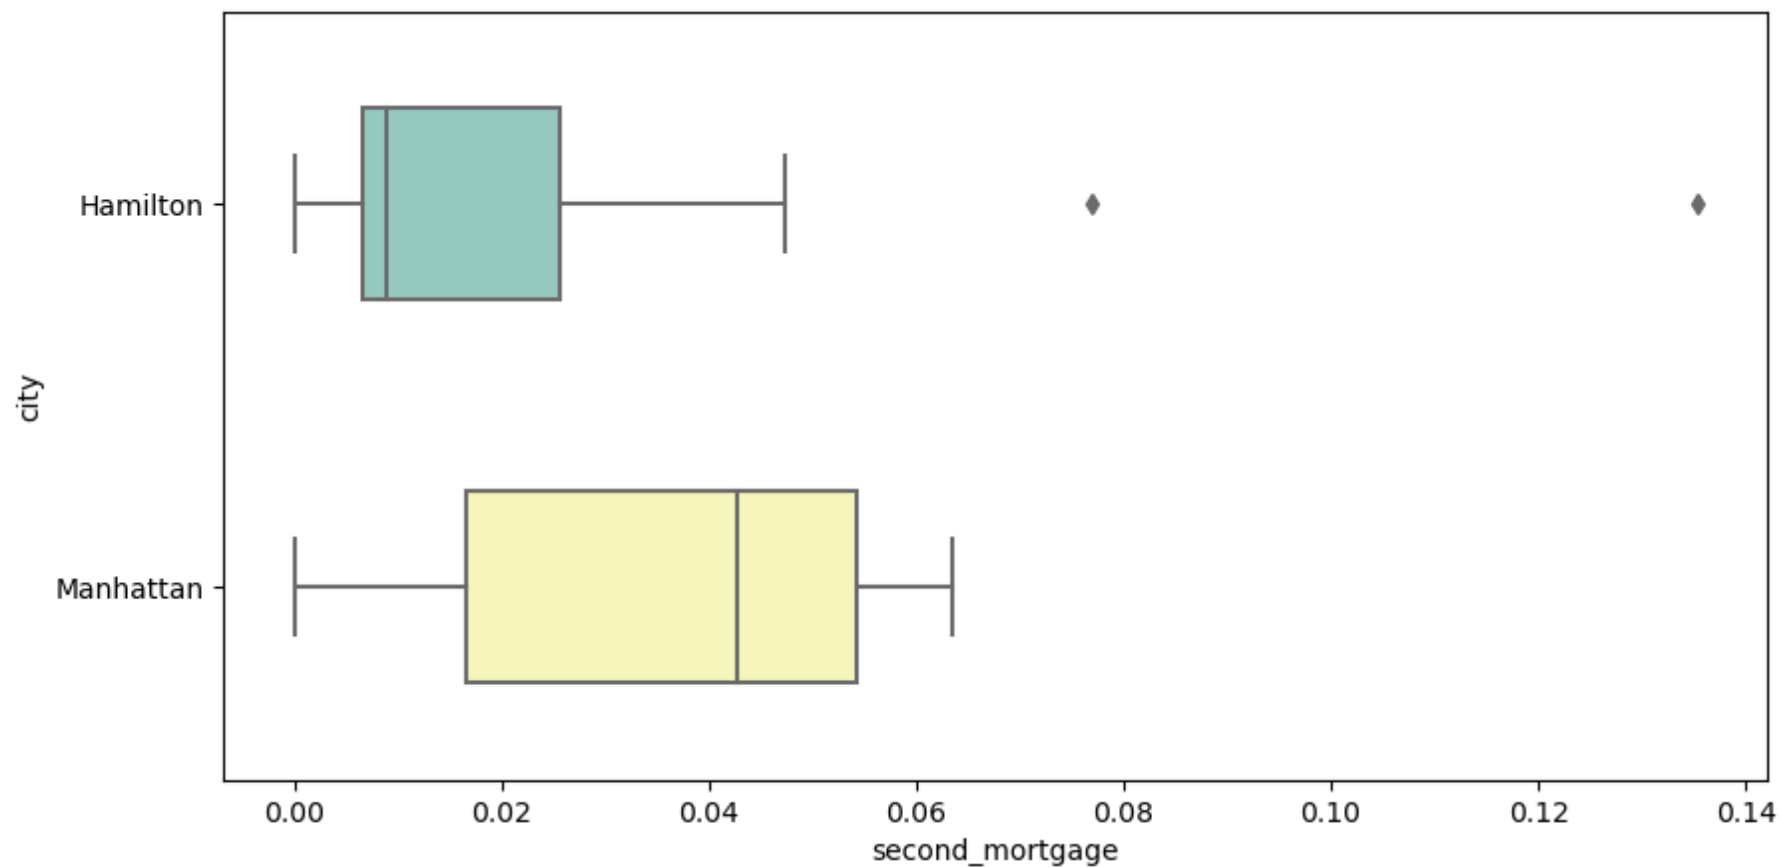
```
Out[46]:
```

	COUNTYID	STATEID	state	state_ab	city	place	type	primary	zip_code	area_code	...	female_age_stdev	female_age_sa
UID													
267822	53	36	New York	NY	Hamilton	Hamilton	City	tract	13346	315	...	22.51276	
263797	21	34	New Jersey	NJ	Hamilton	Yardville	City	tract	8610	609	...	24.05831	
270979	17	39	Ohio	OH	Hamilton	Hamilton City	Village	tract	45015	513	...	22.66500	
259028	95	28	Mississippi	MS	Hamilton	Hamilton	CDP	tract	39746	662	...	22.79602	
270984	17	39	Ohio	OH	Hamilton	New Miami	Village	tract	45013	513	...	24.55724	


5 rows × 79 columns





```
In [47]: plt.figure(figsize=(10,5))  
sns.boxplot(data=df_box_city,x='second_mortgage',y='city',width=0.5,palette='Set3')  
plt.show()
```




```
In [48]: plt.figure(figsize=(10,5))  
sns.boxplot(data=df_box_city,x='home_equity',y='city',width=0.5,palette='Set3')  
plt.show()
```



```
In [49]: plt.figure(figsize=(10,5))  
sns.boxplot(data=df_box_city,x='debt',y='city',width=0.5,palette='Set3')  
plt.show()
```



```
In [50]: plt.figure(figsize=(10,5))  
sns.boxplot(data=df_box_city,x='bad_debt',y='city',width=0.5,palette='Set3')  
plt.show()
```




Manhattan has higher metrics compared to Hamilton

e) Create a collated income distribution chart for family income, house hold income, and remaining income

```
In [51]: sns.distplot(df_train['hi_mean'])  
plt.title('Household Income Distribution Chart')  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



```
In [52]: sns.distplot(df_train['family_mean'])  
plt.title('Family Income Distribution Chart')  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



```
In [53]: sns.distplot(df_train['family_mean']-df_train['hi_mean'])  
plt.title('Remaining Income Distribution Chart')  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



Income Distribution has Normal Distribution

Project Task: Week 2

Exploratory Data Analysis (EDA):

1. Perform EDA and come out with insights into population density and age. You may have to derive new fields (make sure to weight averages for accurate measurements):
 - a) Use pop and ALand variables to create a new field called population density
 - b) Use male_age_median, female_age_median, male_pop, and female_pop to create a new field called median age
 - c) Visualize the findings using appropriate chart type
2. Create bins for population into a new variable by selecting appropriate class interval so that the number of categories don't exceed 5 for the ease of analysis.
 - a) Analyze the married, separated, and divorced population for these population brackets
 - b) Visualize using appropriate chart type
3. Please detail your observations for rent as a percentage of income at an overall level, and for different states.
4. Perform correlation analysis for all the relevant variables by creating a heatmap. Describe your findings.

1. Perform EDA and come out with insights into population density and age. You may have to derive new fields (make sure to weight averages for accurate measurements):


```
In [54]: fig,(ax1,ax2,ax3)=plt.subplots(3,1)
sns.distplot(df_train['pop'],ax=ax1)
sns.distplot(df_train['male_pop'],ax=ax2)
sns.distplot(df_train['female_pop'],ax=ax3)
plt.subplots_adjust(wspace=0.8,hspace=0.8)
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:




`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).





```
In [55]: fig,(ax1,ax2)=plt.subplots(2,1)
sns.distplot(df_train['male_age_mean'],ax=ax1)
sns.distplot(df_train['female_age_mean'],ax=ax2)
plt.subplots_adjust(wspace=0.8,hspace=0.8)
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



a) Use pop and ALand variables to create a new field called population density


```
In [56]: df_train['pop_density'] = df_train['pop']/df_train['ALand']
```

```
In [57]: df_test['pop_density'] = df_test['pop']/df_test['ALand']
```

```
In [58]: sns.distplot(df_train['pop_density'])  
plt.title('Population Density')  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



b) Use male_age_median, female_age_median, male_pop, and female_pop to create a new field called median age

```
In [59]: df_train['age_median'] = (df_train['male_age_median']+df_train['female_age_median'])/2
```

```
In [60]: df_test['age_median'] = (df_test['male_age_median']+df_test['female_age_median'])/2
```

```
In [61]: df_train[['male_age_median', 'female_age_median', 'male_pop', 'female_pop', 'age_median']].head()
```

Out[61]:


	male_age_median	female_age_median	male_pop	female_pop	age_median
UID					
267822	44.00000	45.33333	2612	2618	44.666665
246444	32.00000	37.58333	1349	1284	34.791665
245683	40.83333	42.83333	3643	3238	41.833330
279653	48.91667	50.58333	1141	1559	49.750000
247218	22.41667	21.58333	2586	3051	22.000000

c) Visualize the findings using appropriate chart type

```
In [62]: sns.distplot(df_train['age_median'])  
plt.title('Median Age Distribution')  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



```
In [63]: sns.boxplot(df_train['age_median'])  
plt.title('Population Density')  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning:

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.



2. Create bins for population into a new variable by selecting appropriate class interval so that the number of categories don't exceed 5 for the ease of analysis.


```
In [64]: df_train['pop'].describe()
```

```
Out[64]: count    27321.000000
         mean      4316.032685
         std       2169.226173
         min         0.000000
         25%      2885.000000
         50%      4042.000000
         75%      5430.000000
         max      53812.000000
         Name: pop, dtype: float64
```

```
In [65]: df_train['pop_bins'] = pd.cut(df_train['pop'],bins=5,labels=['very low','low','medium','high','very high'])
```

```
In [66]: df_train[['pop','pop_bins']]
```

```
Out[66]:
```

	pop	pop_bins
UID		
267822	5230	very low
246444	2633	very low
245683	6881	very low
279653	2700	very low
247218	5637	very low
...
279212	1847	very low
277856	4155	very low
233000	2829	very low
287425	11542	low
265371	3726	very low

27321 rows × 2 columns

```
In [67]: df_train['pop_bins'].value_counts()
```

```
Out[67]: very low      27058  
low            246  
medium         9  
high           7  
very high      1  
Name: pop_bins, dtype: int64
```

a) Analyze the married, separated, and divorced population for these population brackets

```
In [68]: df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].count()
```

```
Out[68]:
```

	married	separated	divorced
pop_bins			
very low	27058	27058	27058
low	246	246	246
medium	9	9	9
high	7	7	7
very high	1	1	1

```
In [69]: df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].agg(['mean', 'median'])
```

Out[69]:


	married		separated		divorced	
	mean	median	mean	median	mean	median
pop_bins						
very low	0.507548	0.524680	0.019126	0.013650	0.100504	0.096020
low	0.584894	0.593135	0.015833	0.011195	0.075348	0.070045
medium	0.655737	0.618710	0.005003	0.004120	0.065927	0.064890
high	0.503359	0.335660	0.008141	0.002500	0.039030	0.010320
very high	0.734740	0.734740	0.004050	0.004050	0.030360	0.030360

1. Very high population group has more married people and less percentage of separated and divorced couples
2. In very low population groups, there are more divorced people

b) Visualize using appropriate chart type

```
In [70]: plt.figure(figsize=(10,5))
pop_bin_married = df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].agg('mean')
pop_bin_married.plot(figsize=(20,8))
plt.legend(loc='best')
plt.show()
```

<Figure size 1000x500 with 0 Axes>



3. Please detail your observations for rent as a percentage of income at an overall level, and for different states.

```
In [71]: rent_state_mean = df_train.groupby(by='state')['rent_mean'].agg(['mean'])  
rent_state_mean.head()
```

Out[71]:

mean	
state	
Alabama	774.004927
Alaska	1185.763570
Arizona	1097.753511
Arkansas	720.918575
California	1471.133857

```
In [72]: income_state_mean = df_train.groupby(by='state')['family_mean'].agg(['mean'])  
income_state_mean.head()
```

Out[72]:

mean	
state	
Alabama	67030.064213
Alaska	92136.545109
Arizona	73328.238798
Arkansas	64765.377850
California	87655.470820

```
In [73]: rent_perc_of_income = rent_state_mean['mean']/income_state_mean['mean']  
rent_perc_of_income.head(10)
```

```
Out[73]: state  
Alabama          0.011547  
Alaska           0.012870  
Arizona          0.014970  
Arkansas         0.011131  
California       0.016783  
Colorado         0.013529  
Connecticut      0.012637  
Delaware         0.012929  
District of Columbia 0.013198  
Florida         0.015772  
Name: mean, dtype: float64
```

```
In [74]: sum(df_train['rent_mean'])/sum(df_train['family_mean'])
```

```
Out[74]: 0.013358170721473864
```


4. Perform correlation analysis for all the relevant variables by creating a heatmap. Describe your findings.

```
In [75]: df_train.columns
```

```
Out[75]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',  
              'primary', 'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater',  
              'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',  
              'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',  
              'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',  
              'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',  
              'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',  
              'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',  
              'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',  
              'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',  
              'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',  
              'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',  
              'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',  
              'hs_degree_male', 'hs_degree_female', 'male_age_mean',  
              'male_age_median', 'male_age_stdev', 'male_age_sample_weight',  
              'male_age_samples', 'female_age_mean', 'female_age_median',  
              'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',  
              'pct_own', 'married', 'married_snp', 'separated', 'divorced',  
              'bad_debt', 'bins', 'pop_density', 'age_median', 'pop_bins'],  
              dtype='object')
```

```
In [76]: cor = df_train[['COUNTYID', 'STATEID', 'zip_code', 'pop', 'family_mean', 'second_mortgage', 'home_equity', 'debt',  
                        'hs_degree', 'age_median', 'pct_own', 'married', 'separated', 'divorced']].corr()
```

```
In [77]: plt.figure(figsize=(20,10))
sns.heatmap(cor,annot = True,cmap = 'coolwarm')
plt.show()
```



1. High positive correlation is noticed between pop, male_pop and female_pop
2. High positive correlation is noticed between rent_mean, hi_mean, family_mean, hc_mean

Project Task: Week 3

Data Pre-processing:

1. The economic multivariate data has a significant number of measured variables. The goal is to find where the measured variables depend on a number of smaller unobserved common factors or latent variables. 2. Each variable is assumed to be dependent upon a linear combination of the common factors, and the coefficients are known as loadings. Each measured variable also includes a component due to independent random variability, known as “specific variance” because it is specific to one variable. Obtain the common factors and then plot the loadings. Use factor analysis to find latent variables in our dataset and gain insight into the linear relationships in the data. Following are the list of latent variables:

- Highschool graduation rates
- Median population age
- Second mortgage statistics
- Percent own
- Bad debt expense

```
In [78]: from sklearn.decomposition import FactorAnalysis  
        from factor_analyzer import FactorAnalyzer
```

```
In [79]: fa = FactorAnalyzer(n_factors=5)
fa.fit_transform(df_train.select_dtypes(exclude=('object', 'category')))
fa.loadings_
```

```
Out[79]: array([[ -1.12589165e-01,   1.95646462e-02,  -2.39331065e-02,
        -6.27632576e-02,   4.23474724e-02],
       [ -1.10186762e-01,   1.33506215e-02,   2.79651247e-02,
        -1.49825858e-01,   1.10838804e-01],
       [ -8.28678641e-02,   5.16372369e-02,  -1.36451867e-01,
        -4.98918626e-02,  -1.04024839e-01],
       [  1.80961146e-02,   1.92013750e-02,   5.81329804e-03,
         2.64842729e-02,  -6.12442488e-03],
       [  9.02324755e-02,  -9.72544268e-02,  -6.54601264e-02,
        -1.33145893e-01,  -1.48594590e-01],
       [ -1.07335681e-02,  -4.12376813e-02,   1.45853484e-01,
         8.80433491e-03,   1.08227567e-01],
       [ -4.28796974e-02,  -2.09780212e-02,   3.66726851e-02,
        -9.45597345e-02,   5.91380498e-02],
       [ -2.44243072e-03,  -1.53245405e-02,  -2.68300788e-03,
        -4.52473004e-02,   2.37240637e-02],
       [  7.92164316e-02,   9.57453295e-01,  -8.71151617e-02,
        -6.59924038e-03,  -3.97273194e-02],
       [  7.39808195e-02,   9.18750503e-01,  -1.08834837e-01,
        -2.79371589e-02,  -3.93153647e-02],
       [  8.06598893e-02,   9.47839216e-01,  -6.08006502e-02,
         1.53627080e-02,  -3.86977273e-02],
       [  7.70052110e-01,   9.84675423e-03,  -3.71249731e-02,
         1.14949033e-01,  -1.23784685e-01],
       [  7.18615877e-01,   6.24980466e-03,  -4.59787397e-02,
         1.09109686e-01,  -1.35301910e-01],
       [  7.07647243e-01,   2.46625402e-02,  -1.00860864e-02,
         1.04472486e-01,   7.72381241e-02],
       [ -1.34545492e-01,   3.36809296e-01,  -4.87894961e-01,
        -4.15446193e-02,   3.17608528e-01],
       [  2.31079707e-01,   4.37729793e-01,  -6.40209212e-01,
        -2.52311017e-02,   3.47216227e-01],
       [ -4.52068114e-02,   3.51263837e-02,   3.07537011e-02,
         4.44793494e-01,  -1.63273406e-01],
       [ -2.50717029e-02,   1.70166793e-02,   4.57227084e-02,
         6.76083841e-01,  -1.55256749e-01],
       [ -3.90694436e-02,  -1.67460874e-02,   8.13962691e-02,
         8.36389104e-01,  -9.18259801e-02],
       [ -5.14161948e-02,  -3.57207133e-02,   1.10795166e-01,
         9.25123723e-01,  -4.44866476e-02],
       [ -6.08589985e-02,  -4.41860613e-02,   1.35794025e-01,
```

```
9.53019910e-01, -2.21548664e-02],  
[-4.57771160e-02, -5.25526117e-02, 1.41019874e-01,  
9.32702625e-01, -5.86526030e-07],  
[-4.19486038e-02, -5.90387634e-02, 1.28851786e-01,  
8.87316685e-01, 1.05894282e-02],  
[-2.47894634e-02, -7.29670549e-02, 9.41510484e-02,  
7.79023672e-01, 2.95352803e-02],  
[ 2.12258447e-01, 4.65992339e-01, -6.14495943e-01,  
-2.47660016e-02, 3.66644520e-01],  
[ 2.33057238e-01, 4.47057843e-01, -6.28263418e-01,  
-2.71547710e-02, 3.43419607e-01],  
[ 7.85157086e-01, 4.91249255e-02, 1.44540484e-01,  
-2.05217626e-01, -1.54523360e-01],  
[ 7.10324884e-01, 4.99730438e-02, 1.32239991e-01,  
-2.19171864e-01, -2.10505572e-01],  
[ 8.61780938e-01, 4.35044832e-02, 1.65839096e-01,  
-1.19850813e-01, 3.16733557e-02],  
[-2.23443273e-01, 8.46259553e-01, -4.61177211e-02,  
6.88599244e-02, 2.27742316e-01],  
[ 1.43837555e-01, 9.53197411e-01, 2.27887433e-02,  
-4.57890463e-02, 1.00796445e-01],  
[ 8.30286472e-01, 3.42026001e-02, 1.61105999e-01,  
-2.04570321e-01, -7.48710515e-02],  
[ 7.94476569e-01, 2.83818596e-02, 1.51219547e-01,  
-2.07681490e-01, -9.12497107e-02],  
[ 8.11481647e-01, 4.32314892e-02, 1.43645559e-01,  
-1.07778260e-01, 5.79540142e-02],  
[-3.37741906e-01, 8.64927620e-01, 3.58933696e-02,  
9.07183945e-02, 4.46327252e-02],  
[ 5.03572666e-02, 9.35515342e-01, 1.51475401e-01,  
-2.51501271e-02, -9.34471599e-02],  
[ 9.78242236e-01, -3.31490281e-02, -1.05261173e-01,  
4.50364233e-02, 7.37362022e-02],  
[ 9.59137188e-01, -3.90023001e-02, -1.20630339e-01,  
4.52591414e-02, 6.64877229e-02],  
[ 8.14087168e-01, 2.23057306e-03, 7.66518502e-02,  
2.02747427e-02, 1.27634817e-01],  
[-4.15354001e-01, 7.18339595e-01, 3.40068079e-01,  
-7.18402753e-02, -2.77950528e-01],  
[ 7.64912732e-02, 7.24900618e-01, 2.74193199e-01,  
-4.83952665e-02, -3.52988266e-01],  
[ 9.10390865e-01, -5.36541210e-02, -4.68641894e-02,
```

```
-7.64182289e-04, 1.63870465e-01],  
[ 8.73011863e-01, -5.30302292e-02, -5.89943137e-02,  
-1.58989763e-03, 1.52417542e-01],  
[ 7.55087660e-01, -3.56133725e-03, 5.39542547e-02,  
4.24181451e-03, 2.58043474e-01],  
[-1.23469884e-01, 6.07438127e-01, 6.33039219e-01,  
-2.14798817e-02, 2.47973916e-01],  
[-3.42866892e-01, 5.59526281e-01, 5.88213008e-01,  
-2.51533511e-02, 2.18419885e-01],  
[-1.60867218e-01, -1.53062636e-02, -1.57026580e-01,  
1.09243756e-01, -6.61660830e-01],  
[-1.37306763e-01, -2.17250667e-02, -1.58408927e-01,  
1.25156193e-01, -6.71630803e-01],  
[ 2.45096187e-01, -2.54584596e-02, -2.66691407e-02,  
9.53148453e-02, -6.42510825e-01],  
[ 2.03988663e-01, 7.85172830e-02, -3.01656217e-01,  
2.28379447e-02, -6.29223339e-01],  
[ 1.08926072e-01, -6.34332386e-02, -3.36565230e-02,  
-9.49480417e-02, 6.81473815e-01],  
[-2.63787620e-01, -6.43280757e-03, -3.58792210e-02,  
-9.37962495e-02, 6.47817018e-01],  
[-2.15717048e-01, -7.36588949e-02, 3.50113228e-01,  
-1.95201590e-02, 6.36783755e-01],  
[ 3.94306147e-01, 6.09565684e-02, 2.55337862e-01,  
-2.20362097e-01, -1.84248076e-01],  
[ 4.07877889e-01, 6.27256516e-02, 2.23926907e-01,  
-2.10028735e-01, -1.71989219e-01],  
[ 3.53156876e-01, 5.36715654e-02, 2.69603565e-01,  
-2.16933216e-01, -1.80072062e-01],  
[ 2.33537264e-01, -4.91732958e-02, 8.14450787e-01,  
9.36688926e-02, 3.27131934e-01],  
[ 2.40298202e-01, -3.38140117e-02, 8.31496951e-01,  
7.52417503e-02, 2.46323597e-01],  
[-6.71839476e-02, 6.58504524e-02, 5.86207673e-01,  
8.72955174e-02, 9.12541343e-02],  
[ 5.59835552e-02, 8.17918708e-01, -1.78458350e-01,  
-1.55949439e-02, -3.34299733e-02],  
[ 7.16426403e-02, 9.23428548e-01, -1.07142695e-01,  
-2.78635385e-02, -4.35991115e-02],  
[ 1.92496945e-01, -4.75870400e-02, 8.03173185e-01,  
1.43492708e-01, 3.33862150e-01],  
[ 1.87644433e-01, -3.29941014e-02, 8.58024490e-01,
```

```

1.31329956e-01, 2.55679724e-01],
[-1.02263656e-01, 6.03984253e-02, 4.72982250e-01,
 7.36848367e-02, 1.12273907e-01],
[ 6.14776639e-02, 8.77962739e-01, -1.50410284e-01,
 2.20991026e-02, -4.17158180e-02],
[ 7.83728211e-02, 9.54508776e-01, -5.91095904e-02,
 1.64800910e-02, -4.32590995e-02],
[-3.24381953e-02, 1.11167165e-01, 7.84467411e-01,
 -4.37718523e-02, -2.80931237e-01],
[ 1.76682388e-01, 1.90494238e-01, 5.61405488e-01,
 -1.20746165e-01, -1.32570786e-01],
[-6.37386638e-02, -7.03047917e-02, -2.68934066e-01,
 1.28589791e-01, 1.88507855e-01],
[-1.56051273e-01, -7.08033934e-02, -1.45964500e-01,
 1.24253731e-01, 1.46293109e-01],
[-3.56716294e-01, -5.29910746e-02, 1.47771602e-01,
 2.87196184e-02, 1.13159575e-01],
[ 2.42173831e-01, -2.86199123e-02, -3.25958340e-02,
 1.05027813e-01, -6.55406057e-01],
[ 3.50196741e-01, -1.05016404e-02, -3.95274115e-01,
 5.92876782e-02, 2.91651780e-01],
[ 2.25671545e-01, -3.42672754e-02, 8.92876622e-01,
 1.12426812e-01, 2.67065202e-01]])

```

Project Task: Week 4

Data Modeling :

1. Build a linear Regression model to predict the total monthly expenditure for home mortgages loan. Please refer 'deplotment_RE.xlsx'.
Column hc_mortgage_mean is predicted variable. This is the mean monthly mortgage and owner costs of specified geographical location.
Note: Exclude loans from prediction model which have NaN (Not a Number) values for hc_mortgage_mean.

a) Run a model at a Nation level. If the accuracy levels and R square are not satisfactory proceed to below step.

b) Run another model at State level. There are 52 states in USA.

c) Keep below considerations while building a linear regression model. Data Modeling :

- Variables should have significant impact on predicting Monthly mortgage and owner costs

- Utilize all predictor variable to start with initial hypothesis
 - R square of 60 percent and above should be achieved
 - Ensure Multi-collinearity does not exist in dependent variables
 - Test if predicted variable is normally distributed
1. Build a linear Regression model to predict the total monthly expenditure for home mortgages loan. Please refer 'deplotment_RE.xlsx'.
Column hc_mortgage_mean is predicted variable. This is the mean monthly mortgage and owner costs of specified geographical location.
Note: Exclude loans from prediction model which have NaN (Not a Number) values for hc_mortgage_mean.

```
In [80]: df_train.columns
```

```
Out[80]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
               'primary', 'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater',
               'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',
               'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',
               'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
               'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',
               'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',
               'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
               'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
               'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
               'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
               'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
               'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
               'hs_degree_male', 'hs_degree_female', 'male_age_mean',
               'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
               'male_age_samples', 'female_age_mean', 'female_age_median',
               'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
               'pct_own', 'married', 'married_snp', 'separated', 'divorced',
               'bad_debt', 'bins', 'pop_density', 'age_median', 'pop_bins'],
              dtype='object')
```

```
In [81]: df_train['type'].unique()
type_dict={'type':{'City':1,
                  'Urban':2,
                  'Town':3,
                  'CDP':4,
                  'Village':5,
                  'Borough':6}}
df_train.replace(type_dict,inplace=True)
```

```
In [82]: df_train['type'].unique()
```

```
Out[82]: array([1, 2, 3, 4, 5, 6], dtype=int64)
```

```
In [83]: df_test.replace(type_dict,inplace=True)
```

```
In [84]: df_test['type'].unique()
```

```
Out[84]: array([4, 1, 6, 3, 5, 2], dtype=int64)
```

```
In [85]: feature_cols = ['COUNTYID', 'STATEID', 'zip_code', 'type', 'pop', 'family_mean', 'second_mortgage', 'home_equity', 'debt',
                        'hs_degree', 'age_median', 'pct_own', 'married', 'separated', 'divorced']
```

```
In [86]: x_train = df_train[feature_cols]
y_train = df_train['hc_mortgage_mean']
```

```
In [87]: x_test = df_test[feature_cols]
y_test = df_test['hc_mortgage_mean']
```

```
In [88]: from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error, accuracy_score
```



```
In [89]: x_train.head()
```

```
Out[89]:
```

	COUNTYID	STATEID	zip_code	type	pop	family_mean	second_mortgage	home_equity	debt	hs_degree	age_median	pct_own	ma
UID													
267822	53	36	13346	1	5230	67994.14790	0.02077	0.08919	0.52963	0.89288	44.666665	0.79046	0.5
246444	141	18	46616	1	2633	50670.10337	0.02222	0.04274	0.60855	0.90487	34.791665	0.52483	0.3
245683	63	18	46122	1	6881	95262.51431	0.00000	0.09512	0.73484	0.94288	41.833330	0.85331	0.6
279653	127	72	927	2	2700	56401.68133	0.01086	0.01086	0.52714	0.91500	49.750000	0.65037	0.4
247218	161	20	66502	1	5637	54053.42396	0.05426	0.05426	0.51938	1.00000	22.000000	0.13046	0.1

```
In [90]: sc = StandardScaler()
x_train_scaled = sc.fit_transform(x_train)
x_test_scaled = sc.fit_transform(x_test)
```

a) Run a model at a Nation level. If the accuracy levels and R square are not satisfactory proceed to below step.

```
In [91]: linreg = LinearRegression()
```

```
In [92]: linreg.fit(x_train_scaled,y_train)
```

```
Out[92]: LinearRegression()
```

```
In [93]: y_pred = linreg.predict(x_test_scaled)
```

```
In [94]: print('overall R2 score of linear regression model',r2_score(y_test,y_pred))
```

overall R2 score of linear regression model 0.7348210754610929

```
In [95]: print('Overall RMSE of linear regression model', np.sqrt(mean_squared_error(y_test, y_pred)))
```

Overall RMSE of linear regression model 323.1018894984635

b) Run another model at State level. There are 52 states in USA.

```
In [96]: state = df_train['STATEID'].unique()
```

```
In [97]: state[0:5]
```

```
Out[97]: array([36, 18, 72, 20,  1], dtype=int64)
```

```
In [98]: for i in [20,1,45]:
          print('State ID-',i)

          x_train_nation = df_train[df_train['COUNTYID']==i][feature_cols]
          y_train_nation = df_train[df_train['COUNTYID']==i]['hc_mortgage_mean']

          x_test_nation = df_test[df_test['COUNTYID']==i][feature_cols]
          y_test_nation = df_test[df_test['COUNTYID']==i]['hc_mortgage_mean']

          x_train_scaled_nation = sc.fit_transform(x_train_nation)
          x_test_scaled_nation = sc.fit_transform(x_test_nation)

          linreg.fit(x_train_scaled_nation,y_train_nation)
          y_pred_nation = linreg.predict(x_test_scaled_nation)

          print("Overall R2 score of Linear Regression model for state,",i,":-", r2_score(y_test_nation,y_pred_nation))
          print("Overall RMSE of Linear Regression model for state,",i,":-",np.sqrt(mean_squared_error(y_test_nation,y_pred_nation)))
          print('\n')
```

State ID- 20

Overall R2 score of Linear Regression model for state, 20 :- 0.6046603766461809

Overall RMSE of Linear Regression model for state, 20 :- 307.97188999314716

State ID- 1

Overall R2 score of Linear Regression model for state, 1 :- 0.8104382475484616

Overall RMSE of Linear Regression model for state, 1 :- 307.82758618484354

State ID- 45


Overall R2 score of Linear Regression model for state, 45 :- 0.7887446497855253

Overall RMSE of Linear Regression model for state, 45 :- 225.69615420724134

```
In [99]: residuals = y_test-y_pred  
residuals
```

```
Out[99]: UID  
255504    281.969088  
252676    -69.935775  
276314    190.761969  
248614   -157.290627  
286865    -9.887017  
...  
238088   -67.541646  
242811   -41.578757  
250127  -127.427569  
241096  -330.820475  
287763    217.760642  
Name: hc_mortgage_mean, Length: 11709, dtype: float64
```


```
In [100]: plt.hist(residuals)  
plt.show()
```



```
In [101]: sns.distplot(residuals)  
plt.show()
```


C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



The residuals is Normally Distributed

```
In [102]: plt.scatter(residuals,y_pred)  
plt.show()
```



```
In [ ]:
```